

Logistic Regression for machine Learning

March 10, 2025

```
[24]: ## Loading necessary libraries
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.preprocessing import LabelEncoder
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, confusion_matrix, \
    classification_report, roc_curve, auc
import warnings
warnings.filterwarnings("ignore")
```

```
[25]: ## Import the dataset
df = pd.read_csv("C:/Users/PC/OneDrive/Desktop/Data Science/Datasets/Datasets/
    \Stroke.csv")
```

```
[26]: ## Data Exploration
## View the first view rows of the dataset
df.head()
```

```
[26]:
```

	id	gender	age	hypertension	heart_disease	ever_married	\
0	9046	Male	67.0	0	1	Yes	
1	31112	Male	80.0	0	1	Yes	
2	60182	Female	49.0	0	0	Yes	
3	1665	Female	79.0	1	0	Yes	
4	56669	Male	81.0	0	0	Yes	

	work_type	Residence_type	avg_glucose_level	bmi	smoking_status	\
0	Private	Urban	228.69	36.6	formerly smoked	
1	Private	Rural	105.92	32.5	never smoked	
2	Private	Urban	171.23	34.4	smokes	
3	Self-employed	Rural	174.12	24.0	never smoked	
4	Private	Urban	186.21	29.0	formerly smoked	

	stroke
0	1

```

1      1
2      1
3      1
4      1

```

```
[27]: ## Check the structure of the dataset
df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4909 entries, 0 to 4908
Data columns (total 12 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   id                    4909 non-null   int64
 1   gender                4909 non-null   object
 2   age                   4909 non-null   float64
 3   hypertension          4909 non-null   int64
 4   heart_disease         4909 non-null   int64
 5   ever_married          4909 non-null   object
 6   work_type             4909 non-null   object
 7   Residence_type        4909 non-null   object
 8   avg_glucose_level     4909 non-null   float64
 9   bmi                   4909 non-null   float64
10   smoking_status        4909 non-null   object
11   stroke                4909 non-null   int64
dtypes: float64(3), int64(4), object(5)
memory usage: 460.3+ KB

```

```
[28]: ## Check data types
df.dtypes
```

```

[28]: id                    int64
gender                    object
age                      float64
hypertension              int64
heart_disease             int64
ever_married              object
work_type                 object
Residence_type            object
avg_glucose_level         float64
bmi                      float64
smoking_status            object
stroke                    int64
dtype: object

```

```
[29]: ## Check for duplicates
df.duplicated().sum()
```

```
[29]: 0
```

```
[30]: ## Check for missing values
df.isnull().sum()
```

```
[30]: id                0
gender              0
age                0
hypertension        0
heart_disease        0
ever_married        0
work_type           0
Residence_type      0
avg_glucose_level    0
bmi                 0
smoking_status       0
stroke              0
dtype: int64
```

```
[31]: ## Summary statistics
df.describe()
```


```
[31]:
```

	id	age	hypertension	heart_disease \
count	4909.000000	4909.000000	4909.000000	4909.000000
mean	37064.313506	42.865374	0.091872	0.049501
std	20995.098457	22.555115	0.288875	0.216934
min	77.000000	0.080000	0.000000	0.000000
25%	18605.000000	25.000000	0.000000	0.000000
50%	37608.000000	44.000000	0.000000	0.000000
75%	55220.000000	60.000000	0.000000	0.000000
max	72940.000000	82.000000	1.000000	1.000000

	avg_glucose_level	bmi	stroke
count	4909.000000	4909.000000	4909.000000
mean	105.305150	28.893237	0.042575
std	44.424341	7.854067	0.201917
min	55.120000	10.300000	0.000000
25%	77.070000	23.500000	0.000000
50%	91.680000	28.100000	0.000000
75%	113.570000	33.100000	0.000000
max	271.740000	97.600000	1.000000

```
[32]: ## Data Preprocessing
df["hypertension"] = df["hypertension"].astype("object")
df["heart_disease"] = df["heart_disease"].astype("object")
df["stroke"] = df["stroke"].astype("object")
## Checking for outliers
```

```
numeric_cols = df.select_dtypes(include = "float64")
numeric_cols.boxplot(figsize = (10, 8))
plt.show()
```



```
[33]: ## Remove outliers
      ## Average glucose level
      Q1 = df["avg_glucose_level"].quantile(0.25)
      Q3 = df["avg_glucose_level"].quantile(0.75)
      IQR = Q3 - Q1

      ## Define the lower and upper bound
      lower_bound = Q1 - 1.5 * IQR
      upper_bound = Q3 + 1.5 * IQR

      ## Remove outliers
      df = df[(df["avg_glucose_level"] >= lower_bound) & (df["avg_glucose_level"] <=
        ↪upper_bound)]

      ## bmi
```

```

Q1 = df["bmi"].quantile(0.25)
Q3 = df["bmi"].quantile(0.75)
IQR = Q3 - Q1

## Define the lower and upper bound
lower_bound = Q1 - 1.5 * IQR
upper_bound = Q3 + 1.5 * IQR

## Remove outliers
df = df[(df["bmi"] >= lower_bound) & (df["bmi"] <= upper_bound)]

```

```

[34]: ## One hot encoding
      ## Select categorical columns
      categorical_cols = df.select_dtypes(include = ["object"]).columns
      ## Initialize the label encoder
      label_encoder = LabelEncoder()
      ## Apply label encoding to selected columns
      for col in categorical_cols:
          df[col] = label_encoder.fit_transform(df[col])

```

```

[35]: ## Drop study id
      df = df.drop(columns = ["id"])

```

```

[36]: ## Define Features and Target variable
      X = df.drop(columns = ["stroke"])
      y = df["stroke"]

```

```

[37]: ## Split the dataset into training and testing sets
      X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2,
      ↪ random_state = 42)

```

```

[38]: ## Feature scaling
      scaler = StandardScaler()
      X_train = scaler.fit_transform(X_train)
      X_test = scaler.fit_transform(X_test)

```

```

[50]: ## Train the Logistic Regression Model
      model = LogisticRegression(class_weight = "balanced")
      model.fit(X_train, y_train)

```

```

[50]: LogisticRegression(class_weight='balanced')

```

```

[51]: ## Make Predictions
      y_pred = model.predict(X_test)

```

```


[52]: ## Model Evaluation
      accuracy = accuracy_score(y_test, y_pred)

```

```
print(accuracy)
```

0.7344300822561692

```
[53]: ## Confusion Matrix
conf_matrix = confusion_matrix(y_test, y_pred)
sns.heatmap(conf_matrix, annot = True, fmt = "d", cmap = "Blues")
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.show()
```



```
[54]: ## Classification Report
print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.99	0.73	0.84	817
1	0.11	0.79	0.19	34
accuracy			0.73	851
macro avg	0.55	0.76	0.52	851
weighted avg	0.95	0.73	0.82	851

